

Tutorial 1


ST2137-2420

Material

This tutorial covers material from chapter 1 of the course textbook. It provides practice on basic R functions. There are many ways to code each question, so try them more than once. The following R functions may be helpful: `combn`, `median`, `union`, `intersect`.

Dataset: phones

The dataset `phones` is available from the `MASS` package (which is installed by default with R). It contains two numeric variables, in a list format. Here is a plot of the two variables:



1. Create a dataframe `df1` with the following columns.

```
df1 <- data.frame(x=phones$year, y=phones$calls)
head(df1)
```

```
  x  y
1 50 4.4
2 51 4.7
3 52 4.7
4 53 5.9
5 54 6.6
6 55 7.3
```

2. Write this dataframe to a csv file in the `data/` folder named `phones-2420.csv`.

```
write.csv(df1, "data/phones-2420.csv", row.names=FALSE)
```

These are the first few lines of the file:

```
"x","y"  
50,4.4  
51,4.7  
52,4.7  
53,5.9  
...
```

3. Answer the following queries about the data:

1. How many rows are there in the dataset?
2. How many observations between 100 and 200 million calls?
3. What are the largest 3 and smallest 3 number of calls?
4. In which year was the largest number of calls made?

```
# number of rows:  
NROW(df1)
```

```
[1] 24
```

```
# num. observations between 100 and 200  
sum(df1$y > 100 & df1$y < 200)
```

```
[1] 5
```

```
# smallest 3 calls  
head(sort(df1$y), n=3)
```

```
[1] 4.4 4.7 4.7
```

```
# largest 3 calls  
tail(sort(df1$y), n=3)
```

```
[1] 159 182 212
```

```
# year with largest number of calls.  
df1$x[df1$y == max(df1$y)]
```

```
[1] 69
```

4. In R, matrix multiplication is carried out with the `%*%` operator. For instance, if we have

$$x_{1,2} = \begin{bmatrix} 1 & 1 \end{bmatrix} \quad y_{2,2} = \begin{bmatrix} 0 & 0.5 \\ 0.5 & 0 \end{bmatrix}$$

Then $x \times y$ is computed as

```
x_mat = matrix(c(1,1), nrow=1)  
y_mat = matrix(c(0, 0.5, 0.5, 0), nrow=2)  
# solve(y_mat) # computes inverse of a square matrix  
# t(y_mat)      # returns transpose of a matrix
```

1. Create a 24×2 matrix **X** with the first column all ones, and the second column containing the **year** vector from the **phones** data. Now create a 24×1 matrix **y** containing the **calls** column.
2. Compute the estimate of the slope and intercept for a least-squares best fit to the above data, storing it as **beta_hat**.

$$(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

3. Compute the fitted y-values, storing them as **y_hat**.

$$\mathbf{X}(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

```

X <- cbind(1, df1$x)
y <- matrix(df1$y, ncol=1)
beta_hat <- solve(t(X) %*% X) %*% t(X) %*% y
y_hat <- X %*% beta_hat

```

5. The `lm` function performs the above computation for us in R. Inspect the output object and retrieve the parameters and the fitted values.


```

lm_output <- lm(y ~ x, data=df1)
lm_output$coefficients["x"]

```

x
5.041478

Here is a plot of the line fitted to the data:



6. The fit of the line has been affected by the anomalous points. Here is an algorithm to compute a fitted line that is not so affected by those points:
- Generate all pairwise combinations of observations. (see `combn`)
 - For each pair of points, compute the gradient.
 - Compute the median over all these gradients. This returns the fitted slope.

Write a for-loop that will compute this median. Compare it to the earlier slope. Here is a plot with the new slope in green:


```

all_combn <- combn(24, 2)
all_slopes <- rep(0.0, length=276)
for(ii in 1:276) {
  y_vec <- df1$y[all_combn[,ii]]
  x_vec <- df1$x[all_combn[,ii]]
  all_slopes[ii] <- (y_vec[2] - y_vec[1])/(x_vec[2] - x_vec[1])
}
fit_slope <- median(all_slopes)
fit_intercept <- median(df1$y) - median(df1$x)*fit_slope

## plotting code (not asked for in this question)
plot(x=phones$year, y=phones$calls, col="red", cex=1.2,
     xlab="Year", ylab="Calls (millions)")

```

```
abline(lm_output, col="blue")
abline(a=fit_intercept, b=fit_slope, col="green", lty=2)
```



7. The file `phones.json` contains corrected readings for particular years. The following commands will read the data into R as a list. Replace the data in `df1` at the appropriate years with the corrected call values. *Hint: read up on `match()` function.*

```
library(jsonlite)
corrected_data <- read_json("data/phones.json", TRUE)

matched_rows <- match(corrected_data$year, df1$x)
df1$y[matched_rows] <- corrected_data$corrected_calls
```